

# GCP: Logging, Monitoring, Observability

## 2023 Ivan Vlad S.

### ▼ 1 Monitoring in GCP

#### ▼ Overview

##### ▼ Capture signals

###### ▼ Metrics

- Apps, services, platform, microservices

###### ▼ Logs

- apps, services, platform

###### ▼ Trace

- apps - analyze latency

##### ▼ Visualize and analyze

- Dashboards

- Metrics Explorer

###### ▼ Logs Explorer

- (used to be Logs Viewer)

- Service Monitoring

- Health Checks

- Debugger

- Profiler

##### ▼ manage incidents

- Alerts

- Error Reporting

- SLO

- troubleshoot

#### ▼ Ops-based tools

- monitoring starts with Signal data

- ▼ logging is all about
  - ▼ collect
    - > automatic logging on all app engine, cloud run, GKE and compute engine VMs
  - ▼ analyze
    - > analyze log data in real time with logs explorer, pub/sub, dataflow, and bigquery
    - > analyze archived logs from cloud storage
  - ▼ export
    - > export to cloud storage, or pub/sub, or bigquery
    - > export logs-based metrics to monitoring
  - ▼ retain
    - > data access logs are retained for 1-3650 days, admin logs for 400 days
    - > longer retention available in cloud storage or big query
- ▼ available logs
  - ▼ cloud audit logs
    - who did what, where, admin activity, data access, system event, access transparency
  - ▼ agent logs
    - fluentd agent, common 3P apps, system software
  - ▼ network logs
    - VPC flow, firewall rules, NAT gateway, load balancer
  - error reporting
- ▼ service monitoring
  - understand and troubleshoot intra-service dependencies
- ▼ app performance mgmt tools
  - ▼ debugger
    - real-time app debugging
    - increased collab by sharing debug sessions
    - debug snapshots, logpoints, conditional debugging
  - ▼ trace

- distributed latency analysis
- near real-time
- find performance degradations in apps
- ▼ profiler
  - improve performance and reduce costs
  - understand your apps call patterns
  - low-impact production CPU and heap profiling
- ▼ **2 Avoiding Customer Pain**
  - ▼ Why monitor?
    - monitoring reveals urgent attention, trends, planning, improvements
    - provides continual improvement, dashboards, alerting, debugging
    - Set proper expectations
  - ▼ monitoring systems should address what is broken and why
    - symptom and cause
  - ▼ critical measures
    - ▼ metrics help measure success
      - ▼ business
        - ROI
        - earnings before interest and taxes (EBIT)
        - employer turnover
        - customer churn
      - ▼ software
        - pageviews
        - user registrations
        - click-throughs
        - checkouts
    - ▼ metrics should be SMART
      - specific
      - measureable

- achievable
- relevant
- timebound
- ▼ the 4 golden signals
  - ▼ latency: impacts user experience
    - > indicate emerging issues
      - > may be tied to capacity demands
      - > may be used to show improvements
  - ▼ traffic: indicates current system demand
    - > historical trends are used for capacity planning
      - > core to calculating infra spend
  - ▼ saturation: how full the service is
    - > focuses on most constrained resources
      - > frequently tied to degrading performance
  - ▼ errors: indicates something is failing
    - > may indicate config or capacity issues
      - > can indicate SLO violation
      - > time to alert?
- ▼ SLIs, SLOs, SLAs
  - ▼ SLI = service level indicator (things you measure)
    - quantifiable measure of service availability
  - ▼ SLO = service level objective (an achievable target)
    - a reliability target for an SLI
    - a principled way to agree on the desired reliability of a service
  - ▼ services need SLOs
    - customer happiness test :: happy = meet SLO, not happy = missed SLO
  - ▼ error budgets
    - an SLO implies an acceptable level of unreliability, this is a budget that can be allocated
    - spend on new feature releases and expected system changes, planned downtime, hardware failure, risky experiments
  - ★ most important feature of any system is its reliability

- ▼ choosing a good SLI
  - ▼ needs to be things we can measure that correlate to the happiness of our users
    - ▼ SLI formula:  $(\text{good events} / \text{valid events}) \times 100\%$ 
      - 3-5 SLIs
- ▼ specifying SLIs
  - ▼ request/response
    - availability, latency, quality
  - ▼ data processing
    - coverage, correctness, freshness, throughput
  - ▼ storage
    - throughput, latency
- ▼ developing SLOs and SLIs
  - what performance does the business need?
  - ▼ user expectations are strongly tied to past performance
    - set SLOS based on past performance and business needs
- ▼ **3 Alerting Policies**
  - ▼ Developing an alerting strategy
    - ▼ Alert = automated notification sent by GCP through some notification channel to an external app, ticketing system or human
      - alerts are based on events in a time series
    - ▼ Goal: human gets notified when needed
      - ▼ a service is down, SLOs or SLAs are not being met, something needs to change
        - e.g., When error budget in danger: Alert!
  - ▼ Evaluating alerts
    - Precision (measure of exactness)
    - Recall (measure of completeness)
    - when error count > budget = Alert!
  - ▼ window length
    - window = regular length subdivision of the SLO in total time

- ▼ use small windows
  - faster alert detection, shorter reset time, poor precision
- ▼ use longer windows
  - better precision, reset and detection times longer, spend more error budget before alert
- ▼ add a duration for better precision
  - use multiple conditions for better precision and recall
  - prioritize alerts based on customer impact and SLA
- ▼ Creating alerts
  - ▼ defined using alert policies
    - ▼ alert policy has a name, one or more conditions, notifications, documentation
      - conditions: what's watched and when to alert
      - can control notification channels
    - ▼ use multiple criteria to create resource groups
      - monitor all resources in a group together
- ▼ Creating alerting policies with the CLI
  - both CLI and API require alert policy be defined in JSON file
  - gcloud and the API can create, retrieve, and delete alerting policies
- ▼ Service Monitoring
  - ▼ helps with SLO and alert creation
    - ▼ consolidated services overview
      - error budget details
  - ▼ access through GCP console or Service Monitoring API
    - ▼ select latency or available metrics to act as SLIs
      - compliance periods - set compliance periods, type, and goal
      - configure alert condition for SLO burn rate
    - ▼ use SLIs to easily create SLOs
      - windows-based vs request-based SLOs
    - alerting easily integrated

- ▼ Monitoring critical systems
  - ▼ monitoring is configured via Workspaces
    - ▼ single pane of glass, cross project visibility, monitor resources in GCP and AWS
      - centralize and consolidate resource monitoring
    - ▼ a workspace belongs to a single host project
      - one workspace can monitor multiple projects
      - multiple workspaces can limit access
      - monitor by project for max isolation
    - ▼ IAM roles control user access to workspace
      - monitoring viewer, editor, admin
      - services may. need perms to add metric data: monitoring metric writer
  - ▼ Understanding dashboards
    - ▼ view and analyze metrics
      - ▼ predefined dashboards
        - dashboards broken into charts
  - ▼ Creating charts
    - ▼ start with Metrics Explorer
      - view data that you don't need to display long term on a dashboard
  - ▼ uptime checks
    - ▼ check public service availability
      - ▼ what makes a good uptime check?
        - > protocol, host, and port are appropriate
        - > response checked for specific content

## ▼ **4 Configuring GCP services for observability**

- ▼ Monitoring
  - ▼ OS monitoring agent
    - gathers system and application metrics from VM instances and sends them to monitoring
    - install docs on google site
  - ▼ logging

- ▼ OS logging agent

- streams logs from common 3P apps and system software to GCP logging
- install docs on google site

- ▼ Baking an image

- ▼ goal = org treat the image creation process as a standard DevOps pipeline:
  - > commits to a code base trigger, build jobs, which create, test and deploy images with all requisite software and apps built in, including the logging and monitoring agents
- Base OS install/GCE public image >> hardened OS image >> platform image >> app image
- hashicorp Packer can automate image builds, integrates well with GCP

- ▼ non-vm resources

- ▼ AppEngine

- ▼ standard and flex support logging and monitoring
  - logs viewable under GAE app resource

- ▼ GKE Monitoring & Logging

- K8s Engine Dashboard

- ▼ Prometheus

- ▼ optional monitoring tool for K8s
  - install prometheus and the collector
  - service metrics using Prometheus exposition format can be exported and made visible as external metrics

- ▼ **5 Monitoring Network Security and Audit logs**

- ▼ Network Security & Audit Logs

- ▼ VPC flow logs > part of Andromeda

- ▼ record a sample of network flows
  - record about one out of every 10 packets of network flows sent from and received by the VM instances, including K8s engine nodes
- ▼ enable VPC flow logs per VPC subnet
  - analyze logs in BQ and visualize in Data Studio

- ▼ VPC firewalls

- ▼ firewall rules logging



- enable firewall rule logging in the console
- provide micro-segmentation
- troubleshooting: using rules to catch incorrect traffic
- ▼ Cloud NAT logs
  - ▼ Cloud NAT allows GCE VMs with no external IP to send and receive packets via the internet
    - fully managed service, software defined, grounded in Andromeda
  - ▼ logging allows you to log NAT connections and/or error - TCP and UDP only
    - view filtering logs in Logs Explorer
- ▼ Packet mirroring
  - ▼ visualize and protect your network, clones VPC instance traffic and FWDs for examination
    - happens at NIC, not part of VPC
- ▼ network intelligence center
  - ▼ centralized network monitoring visibility
    - > Topology: view VPC topology and metrics
    - > Connectivity tests: prevent outages
    - > Performance dashboard: packet loss metrics aggregated across zones
    - > Firewall insights: metrics help understand and optimize firewall configs
- ▼ Audit logs
  - ▼ who did what, where, and when?
    - >> admin activity >> data access >> system event
  - data access logs need to be enabled, not enabled by default

## ▼ 6 Investigating App Performance Issues

- ▼ Debugger
  - ▼ inspect state of a running app in real-time without stopping or slowing it down
    - debugger must be enabled
    - dynamically add log messages with Log Points
- ▼ Trace

- ▼ Cloud Trace tracks app latency
  - trace = collection of spans, span = object that wraps metrics and other contextual info about a unit of work in your app
- ▼ Profiler
  - ▼ statistical, low-overhead memory and CPU profiler > understand performance
  - ▼ CPU time, Heap, allocated heap, contention, threads, Wall time (how long it takes to run a block of code)
    - Subtopic 1